

CS294 Bibliography

This bibliography will be extended as the course progresses.

1 Introduction, overview, terminology

Much of the introductory terminology is taken from the book *Virtual Machines* by James E. Smith and Ravi Nair, Morgan Kaufman, 2005. This book is an excellent introduction to the field, with a superb historical and scholarly overview (but is now 10 years old and could use a refresh).

The *Java Virtual Machine Specification* is available as a download from <http://docs.oracle.com/javase/specs/>. We will be looking at the Java SE 8 JVM Specification in great detail in a coming lecture and you should download the spec. for reference.

VirtualBox is available from virtualbox.org. There is documentation and source code, but I am unaware of any technical overview of the internals.

Technical information on *Rosetta* is hard to come by; Google and Wikipedia are good starting points.

Dynamo: a transparent dynamic optimization system

Vasanth Bala, Evelyn Duesterwald, Sanjeev Banerjia, POPL 2001

<http://dx.doi.org/10.1145/358438.349303>

A widely-cited trace-driven binary reoptimizer. Wikipedia (Tracing_just-in-time_compilation) claims it was the first but I believe that honor should go to *Wiggins-Redstone*, done at DEC (it is hard to find descriptions in the literature; there is a presentation at HotChips 11 [1999]). *Shade* (Cmelik and Keppel, 1993), cited in the *Dynamo* paper, incorporated many of the same ideas but for a different purpose, and that paper cites many earlier systems with elements of tracing; the core idea is very old.

The Transmeta Code Morphing™ Software: using speculation, recovery, and adaptive retranslation to address real-life challenges

James C. Dehnert, Brian K. Grant, John P. Banning, Richard Johnson, Thomas Kistler, Alexander Klaiber, Jim Mattson, CGO '03, pp.15—24.

The best overview of the Transmeta binary translator.

Transmeta Breaks x86 Low-power Barrier

Tom R. Halfhill, Microprocessor Report, Feb. 14, 2000

2. Execution mechanisms, Part I: Interpretation

The origins of AST interpretation seem to be lost in the mists of time — I cannot find anything that is plausibly an original reference.

Slim binaries

Michael Franz, Thomas Kistler, CACM 40(12), Dec 1997

<http://dx.doi.org/10.1145/265563.265576>

A distribution format for programs in the form of compressed ASTs.

3 Anatomy of a Virtual Machine

See section 1 for a reference to the JVM Spec.

Another excellent description of a virtual machine can be found in *Smalltalk-80: The Language and its Implementation*, by Adele Goldberg and Dave Robson, Addison-Wesley, 1983. A free downloadable version is available from the ACM Classic Books web site, <http://dl.acm.org/classics.cfm>. The last section of the book has an executable specification of the Smalltalk-80 VM, written in Smalltalk.

Java Class Viewer is available from <http://www.codeproject.com/Articles/35915/Java-Class-Viewer>. I found that when I unzipped the archive, I ended up with filenames with DOS-like embedded paths. After editing those out of the names, I ran the utility like this (all on one line):
`java -cp JavaClassViewer.jar:CommonLib.jar:FormatCLASS.jar \`
`org.freeinternals.javaclassviewer.Main`
YMMV.

The JNI spec is available at <https://docs.oracle.com/javase/8/docs/technotes/guides/jni/spec/jniTOC.html>.

4 Bytecode interpretation

An overview of the combined area of abstract and virtual machines can be found in:

Abstract machines for programming language implementation

Stephan Diehl, Pieter Hartel, Peter Sestoft, FGCS 16 (2000)

http://www.inf.ed.ac.uk/teaching/courses/lsi/diehl_abstract_machines.pdf

Branch Prediction and the Performance of Interpreters - Don't Trust Folklore

Erven Rohou, Bharath Narasimha Swamy, André Seznec, CGO 2015

<https://hal.inria.fr/hal-01100647>

A recent paper measuring interpreter branch prediction performance on modern hardware.

5 JVM bytecodes

Java intermediate bytecodes

James Gosling, ACM SIGPLAN IR'95 workshop on intermediate representations, 1995

<http://dx.doi.org/10.1145/202530.202541>

6 Dynamic language VMs

Smalltalk

Efficient Implementation of the Smalltalk-80 System

L Peter Deutsch, Allan M Schiffman

POPL 1984, pp 297—302.

<http://dx.doi.org/10.1145/800017.800542>

Some background on Smalltalk-80:

The August 1981 issue of *Byte* magazine was a special issue on the Smalltalk-80 system. Within can be found many articles about the language, environment and implementation.

<https://archive.org/details/byte-magazine-1981-08>

The language, the VM and the interactive system are described in two books:

Smalltalk-80: The language and its implementation

Adele Goldberg and David Robson, Addison-Wesley, 1983.

(aka “the Blue Book”)

Available free from at www.acm.org/classics

Smalltalk-80: The Interactive Programming Environment

Adele Goldberg, Addison-Wesley, 1983

(aka “the Orange Book”)

Available from <http://www.world.st/learn/books>

An additional book collects together implementor’s early experiences:

Smalltalk-80: Bits of history, words of advice

Glenn Krasner (ed), Addison-Wesley, 1983.

(aka “the Green Book”)

Available from <http://www.world.st/learn/books>

The Design and Evaluation of a High Performance Smalltalk System

David Michael Ungar, MIT Press, 1986

Ungar’s UCB thesis presented and evaluated the hardware and software design of a Smalltalk system on a RISC (SOAR). Of its many contributions, the one with most lasting impact has been the Generation Scavenging automatic storage reclamation technique, still widely used (with some enhancements). Various spin-out papers from the SOAR project described aspects of the work:

Compiling Smalltalk-80 to a RISC <http://dx.doi.org/10.1145/36206.36192>

SOAR: Smalltalk without bytcodes <http://dx.doi.org/10.1145/28697.28708>

Architecture of SOAR: Smalltalk on a RISC <http://dx.doi.org/10.1145/800015.808182>

What Price Smalltalk? <http://dx.doi.org/10.1109/MC.1987.1663359>

Generation Scavenging: A non-disruptive high performance storage reclamation algorithm <http://dx.doi.org/10.1145/800020.808261> (this will be covered in the section on GC)

Self Language

Self: The Power of Simplicity

David Ungar and Randall B. Smith, OOPSLA 1987

<http://dx.doi.org/10.1145/38765.38828>

A gentle introduction to the language and the abstract model of computation.

The “Self videos” are about 1h10m and 20m long and also (somewhat immodestly) recommended:

<http://www.selflanguage.org/>

Self VM internal organization

Object storage and inheritance for Self, a prototype-based object-oriented programming language

Elgin Lee, Stanford Engineer’s Thesis, 1988.

The best description of the Self object storage system. I will upload a copy of this to the resources section; it’s hard to get otherwise.

An efficient implementation of SELF, a dynamically-typed object-oriented language based on prototypes

Craig Chambers, David Ungar, Elgin Lee, 1991

<http://dx.doi.org/10.1145/74878.74884>

The second best reference on the Self object storage system, but much less detailed than Lee’s thesis. Also includes detail in the early Self dynamic compiler.

7 Memory management

There are two excellent textbooks covering garbage collection:

Garbage Collection

Richard Jones with Rafael Lins, Wiley 1999 (2nd printing).

The Garbage Collection Handbook

Richard Jones, Antony Hosking, Eliot Moss, CRC Press 2012.

The latter is an updated and extended version of the former, and is preferable, if you have to pick one. The former has more detail on some techniques of historical interest.

In the GC Handbook, Chapter 1 covers mark-sweep, Chapter 3 covers mark-compact, Chapter 4 covers copying collection, Chapter 5 covers reference counting and Chapter 7 cover allocation. Richard Jones maintains a web site about memory management with a comprehensive bibliography: <http://www.cs.kent.ac.uk/people/staff/rej/gc.html>

The assigned reading is *Generation Scavenging: A non-disruptive high performance storage reclamation algorithm*, David Ungar, 1984.

<http://dx.doi.org/10.1145/800020.808261>

The pauseless GC algorithm

Cliff Click, Gil Tene and Michael Wolf, VEE '05

<http://dx.doi.org/10.1145/1064979.1064988>

MMU was introduced in:

A parallel, real-time garbage collector

Perry Cheng and Guy Blelloch, PLDI 2001

<http://dx.doi.org/10.1145/381694.378823>

BMU was introduced in:

MC2: High-Performance Garbage Collection for Memory-Constrained Environments

Narendran Sachindran J. Eliot B. Moss Emery D. Berger, OOPSLA 2004

<http://dx.doi.org/10.1145/1035292.1028984>

Here are a couple of recent papers on the energy impact of GC:

Impact of GC design on power and performance for Android

Ahmed Hussein, Mathias Payer, Antony Hosking and Christopher A Vick, SYSTOR '15

<http://dx.doi.org/10.1145/2757667.2757674>

Don't race the memory bus: taming the GC leadfoot

Ahmed Hussein, Antony Hosking, Mathias Payer and Christopher A Vick, ISMM 2015

<http://dx.doi.org/10.1145/2754169.2754182>

8 Advanced interpretation

Threaded code

James R Bell, CACM 1973 16(6)

<http://dx.doi.org/10.1145/362248.362270>

Indirect threaded code

Robert Dewar, CACM 1975 18(6)

<http://dx.doi.org/10.1145/360825.360849>

Stack caching for interpreters

M Anton Ertl, PLDI 95

<http://dx.doi.org/10.1145/207110.207165>

Combining stack caching with dynamic super instructions

M Anton Ertl and David Gregg, IVME '04

<http://dx.doi.org/10.1145/1059579.1059583>

The Structure and Performance of Efficient Interpreters

M Anton Ertl and David Gregg

<http://www.complang.tuwien.ac.at/anton/tmp/interpreter-arch.ps>

Optimizing direct threaded code by selective inlining

Ian Piumarta and Fabio Riccardi, PLDI 98

<http://dx.doi.org/10.1145/277652.277743>

vmgen — A Generator of Efficient Virtual Machine Interpreters

M Anton Ertl, David Gregg, Andreas Krall, Bernd Paysan, SP&E 2002 32(3)

<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.98.873&rep=rep1&type=pdf>

9 Dynamic compilation

A Brief History of Just-In-Time

John Aycock, Computing Survey, 2003

<http://dx.doi.org/10.1145/857076.857077>

A survey of the dynamic compilation literature.

Efficient Implementation of the Smalltalk-80 System

L Peter Deutsch, Allan M Schiffman

POPL 1984, pp 297—302.

<http://dx.doi.org/10.1145/800017.800542>

Described the first JIT compiler and inline caches.

Self VM — compilation-related

Optimizing Dynamically-Typed Object-Oriented Languages With Polymorphic Inline Caches

Urs Hölzle, Craig Chambers and David Ungar, ECOOP 91

<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.36.6379&rep=rep1&type=pdf>

A generalization of inline caching, paving the way for adaptive optimization.

Customization: optimizing compiler technology for SELF, a dynamically-typed object-oriented programming language

Craig Chambers and David Ungar, PLDI '89

<http://dx.doi.org/10.1145/74818.74831>

<http://dx.doi.org/10.1145/989393.989425> (including a retrospective)

Debugging optimized code with dynamic deoptimization

Urs Hölzle, Craig Chambers and David Ungar, PLDI '92

<http://dx.doi.org/10.1145/143103.143114>

The canonical reference on deoptimization.

Adaptive optimization for Self: Reconciling High Performance with Exploratory Programming

Urs Hölzle, Stanford Ph.D. thesis and Sun Labs technical report, 1994

<http://bibliography.selflanguage.org/urs-thesis.html>

The best complete overview of the last generation of the Self VM. It is summarized in:

A Third-Generation Self Implementation: Reconciling Responsiveness with Performance

Urs Hölzle and David Ungar, OOPSLA 94

<http://dx.doi.org/10.1145/191081.191116>

A complete list of Self publications is at bibliography.selflanguage.org.

HotSpot JVM

A simple graph-based intermediate representation

Cliff Click and Michael Paleczny

<http://dx.doi.org/10.1145/202530.202534>

The Java HotSpot Server Compiler

Michael Paleczny, Chris Vick and Cliff Click, JVM'01

https://www.usenix.org/legacy/events/jvm01/full_papers/paleczny/paleczny.pdf

Fast subtype checking in the HotSpot JVM

Cliff Click and John Rose, Java Grande 2002

<http://dx.doi.org/10.1145/583810.583821>

Design of the Java HotSpot™ client compiler for Java 6

Kotzmann et al, TACO 5(1) 2008

<http://dx.doi.org/10.1145/1369396.1370017>

Miscellaneous compilation and optimization techniques

Optimization of Object-Oriented Programs Using Static Class Hierarchy Analysis

Jeffrey Dean, David Grove, and Craig Chambers, ECOOP 95

<http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.117.2420>

Fast static analysis of C++ virtual function calls

David F Bacon and Peter F Sweeney, OOPSLA 96

<http://dx.doi.org/10.1145/236337.236371>

Introduces Rapid Type Analysis, another technique for devirtualizing calls.

Design, implementation and evaluation of adaptive recompilation with on-stack replacement

Stephen J Fink and Feng Qian, CGO '03

[http://dl.acm.org/citation.cfm?](http://dl.acm.org/citation.cfm?id=776261.776288&coll=DL&dl=GUIDE&CFID=545999259&CFTOKEN=65787538)

[id=776261.776288&coll=DL&dl=GUIDE&CFID=545999259&CFTOKEN=65787538](http://dl.acm.org/citation.cfm?id=776261.776288&coll=DL&dl=GUIDE&CFID=545999259&CFTOKEN=65787538)

Partial Escape Analysis and Scalar Replacement for Java

Lukas Stadler, Thomas Würthinger and Hanspeter Mössenböck, CGO 14

<http://dx.doi.org/10.1145/2544137.2544157>

Automatic Construction of Inlining Heuristics using Machine Learning

Sameer Kulkarni, John Cavazos, Christian Wimmer and Douglas Simon, CGO 2013

<http://dx.doi.org/10.1109/CGO.2013.6495004>

Speculation without regret: reducing deoptimization meta-data in the Graal compiler

Gilles Duboscq, Thomas Würthinger and Hanspeter Mössenböck, PPPJ '14

<http://dx.doi.org/10.1145/2647508.2647521>

Tracing and meta-tracing

See the Dynamo paper (above, under *Introduction*) for an earlier description of trace compilation.

HotpathVM: an effective JIT compiler for resource-constrained devices

Andreas Gal, Christian Probst and Michael Franz, VEE '06

<http://dx.doi.org/10.1145/1134760.1134780>

PyPy's approach to virtual machine construction

Armin Rigo and Samuel Pedroni, OOPSLA '06

<http://dx.doi.org/10.1145/1176617.1176753>

Tracing the meta-level: PyPy's tracing JIT compiler

Carl Friedrich Bolz, Antonio Cuni, Maciej Fijalkowski and Armin Rigo, ICPOOLPS '09

<http://dx.doi.org/10.1145/1565824.1565827>

Runtime feedback in a meta-tracing JIT for efficient dynamic languages

Carl Friedrich Bolz et al., ICPOOLPS 11

<http://dx.doi.org/10.1145/2069172.2069181>

10 Metacircular VMs

See *Smalltalk-80: The language and its implementation* (under the Smalltalk section, above) for a metacircular Smalltalk bytecode interpreter.

Implementing a Java Virtual Machine in the Java Programming Language

Antero Taivalsaari, Sun Microsystems Laboratories Technical Report TR-98- 64, March 1998.

<http://dl.acm.org/citation.cfm?id=974968>

A Java bytecode interpreter written in Java.

The Wikipedia pages referenced in the lectures mention dozens of language implementations stacked on top of the JVM and CLI. Here a couple of earlier such stacked languages, on top of Self:

self includes: Smalltalk

Mario Wolczko, in *Prototype-Based Programming*, Noble, Taivalsaari, Moore (eds), Springer, 1999.

<http://merlinter.com/download/mario.pdf>

Describes an implementation of Smalltalk-80 atop the Self VM.

Design and implementation of Pep, a Java just-in-time translator

Ole Agesen, *Theory and Practice of Object Systems*, 3(2), 1997

<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.22.1089&rep=rep1&type=ps>

An implementation of the JVM atop the Self VM.

Back to the future: the story of Squeak, a practical Smalltalk written in itself
Dan Ingalls, Ted Kaehler, John Maloney, Scott Wallace, Alan Kay, OOPSLA 97
<http://dx.doi.org/10.1145/263700.263754>

SqueakJS: A modern and practical Smalltalk that runs in any browser
Bert Freudenberg et al., DLS 14
<http://dx.doi.org/10.1145/2775052.2661100>
Squeak on JavaScript!

A Java virtual machine architecture for very small devices
Nik Shaylor, Douglas N. Simon, William R Bush, LCTES 03
<http://dx.doi.org/10.1145/780731.780738>
The Squawk JVM, a small JVM implemented in Java and translated to C, inspired by Squeak

The Jalapeño Dynamic Optimizing Compiler for Java
Burke et al., Java Grande 1999
<http://dx.doi.org/10.1145/304065.304113>
The first paper about the system which would eventually be named Jikes RVM.

Adaptive optimization in the Jalapeño JVM
Matthew Arnold, Stephen Fink, David Grove, Michael Hind and Peter Sweeney, OOPSLA 2000
<http://dx.doi.org/10.1145/354222.353175>

Demystifying Magic: High-level Low-level Programming
Frampton et al., VEE '09
<http://dx.doi.org/10.1145/1508293.1508305>
An overview of the techniques used to develop low-level code (VMs, GC) in a HLL (Java)

Oil and Water? High Performance Garbage Collection in Java with MMTk
Stephen M. Blackburn, Perry Cheng, Kathryn M. McKinley, ICSE '04
<http://dl.acm.org/citation.cfm?id=998675.999420&coll=DL&dl=GUIDE&CFID=545999259&CFTOKEN=65787538>

Constructing a metacircular virtual machine in an exploratory programming environment
David Ungar, Adam Spitz and Alex Ausch, OOPSLA '05
<http://dx.doi.org/10.1145/1094855.1094865>
Describes the Klein VM.

Maxine: An approachable virtual machine for, and in, Java
Christian Wimmer et al., TACO 9(4) 2013
<http://dx.doi.org/10.1145/2400682.2400689>

Snippets: Taking the High Road to a Low Level
Doug Simon et al., TACO 12(2) 2015
<http://dx.doi.org/10.1145/2764907>

Truffle/Graal and related topics

Partial Evaluation of Computation Process – An Approach to a Compiler-Compiler

Yoshihiko Futamura, Higher-Order and Symbolic Computation 12, 381–391 (1999)

<http://www.brics.dk/~hosc/local/HOSC-12-4-pp381-391.pdf>

The original description of partial evaluation. Don't be fooled by the '99 publication date — the original (in Japanese) was from the 1970s. See Futamura's web site: <http://fi.ftmr.info/>

One VM to rule them all

Wuerthinger et al., Onward! 2013

<http://dx.doi.org/10.1145/2509578.2509581>

Truffle and Graal papers and presentations can be found at

<https://wiki.openjdk.java.net/display/Graal/Publications+and+Presentations>

Truffle tutorial video at SPLASH 2014, Christian Wimmer

[Video](#) (90mins, SimpleLanguage begins at 22:25) [Slides](#)

Truffle/Graal tutorial at DSLDI 2015, Thomas Würthinger, Christian Humer

Videos: [Part1](#) [Part2](#) (2x90mins)

The first hour of Part 2 is about Truffle. Part 1 is more about Graal; jump at 25mins — it will show the low-level detail of Graal code generation for SimpleLanguage, and includes an introduction to the Ideal Graph Visualizer.

Self-Optimizing AST Interpreters

Thomas Würthinger, Andreas Wöß, Lukas Stadler, Gilles Duboscq, Doug Simon, Christian Wimmer, DLS '12

A Domain-Specific Language for Building Self-Optimizing AST Interpreters

Christian Humer, Christian Wimmer, Christian Wirth, Andreas Wöß, Thomas Würthinger, GPCE'14

An overview of the Truffle DSL. The current version has been enhanced since the paper. For details, see <http://mail.openjdk.java.net/pipermail/graal-dev/2015-February/002912.html>. The change log is [here](#).

Truffle tutorial videos at CGO 2014, Christian Wimmer

Videos: [Part 1](#), [Part 2](#) (2x90mins) [Slides](#)

An Object Storage Model for the Truffle Language Implementation Framework

Andreas Wöß, Christian Wirth, Daniele Bonetta, Chris Seaton, Christian Humer and Hanspeter Mössenböck, PPPJ '14, <http://dx.doi.org/10.1145/2647508.2647517>

Dynamically composing languages in a modular way: supporting C extensions for dynamic languages

Matthias Grimmer, Chris Seeton, Thomas Würthinger and Hanspeter Mössenböck, MODULARITY 2015, <http://dx.doi.org/10.1145/2724525.2728790>

Concurrency

Eliminating synchronization-related atomic operations with biased locking and bulk rebiasing

Kenneth Russell and David Detlefs, OOPSLA '06

<http://dx.doi.org/10.1145/1167515.1167496>

See Jones, Hosking and Moss (*GC Handbook*, above under Memory Management) for descriptions and references on concurrent and parallel GC.

Tools

Building Debuggers and Other Tools: We Can “Have it All”: A Position Paper

Michael L. Van De Vanter, IC00OLPS 2015

<http://vandevanter.net/mlvdv/publications/2015-icoolps.pdf>

Debugging at Full Speed

Chris Seaton, Michael L. Van De Vanter, and Michael Haupt, Dyla'14

<http://dx.doi.org/10.1145/2617548.2617550>

System VMs

Formal requirements for virtualizable third generation architectures

Gerald J. Popek and Robert P. Goldberg, CACM: Volume 17 Issue 7, July 1974

<http://dx.doi.org/10.1145/361011.361073>

The Evolution of an x86 Virtual Machine Monitor

Ole Agesen, Alex Garthwaite, Jeffrey Sheldon, Pratap Subrahmanyam, SIGOPS 44(4), 2010.

<http://dx.doi.org/10.1145/1899928.1899930>

See also the Smith and Nair book (first entry under Introduction, above).